

Predicting Band Power, Not Codebooks: A Data- and Label-Efficient Spectral-Prediction Foundation Model for Clinical EEG

Anonymous submission

Abstract

Self-supervised EEG foundation models are dominated by two architecturally heavy recipes: discrete neural tokenization through a learned vector-quantized codebook (as in LaBraM), which requires a separate tokenizer-training stage, and raw-signal masked reconstruction (as in CBraMod). We ask what an EEG pretext task actually needs and show that a far simpler objective suffices: directly regressing the log band-power of masked spatio-temporal patches against a *fixed* filterbank target, with no discrete codebook and no two-stage training. Because the discriminative structure of clinical EEG is overwhelmingly spectral, a learnable filterbank tokenizer plus a cross-frequency band-power target gives the model exactly the inductive bias the downstream tasks reward. Pretrained on only $\sim 2,500$ hours of EEG—a random one-seventh subset of the TUH EEG Corpus, comparable to the smallest prior foundation model and roughly a quarter of CBraMod’s data—our model reaches state-of-the-art sleep staging (IS-RUC, HMC) and stays competitive across clinical spectral tasks. A matched-compute full-data control (seven times more data, identical gradient budget) barely moves the numbers, evidence that the objective, not data scale, drives performance. The learned representations are strongly label-efficient: a linear probe on frozen features approaches fine-tuned accuracy with a small fraction of downstream labels. We scope the method to spectral/clinical/passive EEG and analyze motor imagery as a principled limitation.

Introduction

Large models pretrained on the TUH EEG Corpus (TUEG) have become the standard recipe for EEG representation learning (Jiang, Zhao, and Lu 2024; Wang et al. 2025). The two dominant designs are both heavy. LaBraM (Jiang, Zhao, and Lu 2024) first trains a vector-quantized neural tokenizer and then predicts its discrete codes, a two-stage pipeline whose first stage is itself a nontrivial spectral autoencoder. CBraMod (Wang et al. 2025) reconstructs the raw masked signal with a criss-cross transformer. Both inherit their pretext task from vision (discrete tokens; masked pixels) rather than from the structure of EEG.

This is an anonymized submission for review purposes only. Distribution, citation, or public sharing of this manuscript is strictly prohibited. Copyright and publication details will appear in the final version if accepted.

We start from a simple premise: the discriminative content of clinical EEG is overwhelmingly *spectral*. Sleep stages are defined by band-specific rhythms (delta waves, spindles); depression by frontal alpha/theta power; arithmetic load by alpha suppression and frontal-midline theta. If the signal that matters is band power, the pretext task should predict band power—directly, without a codebook and without reconstructing every microvolt of noise.

We introduce a **codebook-free spectral-prediction** foundation model. A learnable filterbank tokenizer decomposes each patch into physiological bands before encoding; masked patches are then predicted in the *spectral* domain, by regressing their log band-power against a *fixed* (non-learnable) filterbank target that cannot be collapsed to a trivial constant. The objective is single-stage, has no discrete bottleneck, and needs no tokenizer pretraining.

Beyond simplicity, our central empirical finding is about *efficiency*. Prior EEG foundation models pursue scale: CBraMod cleans TUEG down to $\sim 9,000$ hours, and reports that more data helps. We find the opposite for spectral tasks. Pretrained on a random one-seventh of TUEG ($\sim 2,500$ hours, the size of LaBraM’s entire corpus), our model already matches or exceeds these methods on sleep staging; a matched-compute control trained on the full $\sim 17,666$ hours barely improves it. Data scale is not the bottleneck; the objective is.

Our contributions:

- **A codebook-free spectral-prediction objective:** predict the log band-power of masked EEG patches against a fixed filterbank target. Single-stage, no vector quantization, no reconstruction of raw noise. A fixed target is necessary—a learnable target collapses.
- **Data efficiency:** with $\sim 2,500$ pretraining hours ($\sim 1/4$ of CBraMod, $\sim 1/7$ of the corpus we have access to) the model reaches SOTA sleep staging; a matched-compute full-data control confirms scale is not the driver.
- **Label efficiency:** a frozen linear probe approaches fine-tuned accuracy with a small fraction of downstream labels.
- **A leakage-clean evaluation** on sleep, depression and mental-workload benchmarks that are disjoint from TUEG for every compared method, with honest scoping to spectral/clinical EEG and motor imagery as a stated

limitation.

Related Work

Task-specific EEG decoders. Before foundation models, EEG decoding was dominated by compact task-specific networks: convolutional designs such as EEGNet (Lawhern et al. 2018) and SPaRCNet (Jing et al. 2023), transformer hybrids such as EEG Conformer (Song et al. 2023), CNN-Transformer (Peh, Yao, and Dauwels 2022) and ST-Transformer (Song et al. 2021), and self-supervised or contrastive encoders such as ContraWR (Yang et al. 2023) and FFCL (Li et al. 2022). These are trained per task and do not transfer a shared representation across datasets; we compare against all of them on the identical splits.

What EEG foundation models predict. Recent models pre-train a single encoder on the Temple University Hospital EEG corpus (TUEG) (Obeid and Picone 2016) and differ chiefly in *what the masked-prediction target is*. Two families dominate, and both predict a *proxy* of the signal rather than the discriminative content itself. (i) *Discrete codes*. LaBraM (Jiang, Zhao, and Lu 2024) first trains a vector-quantized neural tokenizer—itsself a spectral autoencoder over Fourier amplitude and phase—and then predicts its discrete indices, a two-stage pipeline in the spirit of BEiT (Bao et al. 2022) and VQ-VAE (van den Oord, Vinyals, and Kavukcuoglu 2017); the codebook is an extra module to train, size and tune. (ii) *Raw-signal reconstruction*. CBraMod (Wang et al. 2025) removes the codebook but reconstructs the *original raw waveform* of masked patches with a criss-cross transformer, following the masked-autoencoder recipe (He et al. 2022); CSBrain (Chen et al. 2025) keeps this raw-waveform target and adds cross-scale spatio-temporal modeling; BIOT (Yang, Westover, and Sun 2023) likewise tokenizes and reconstructs biosignals. Reconstructing the waveform spends model capacity on microvolt-level noise the downstream task never uses. Our target is neither a code nor the waveform: we predict the *log band-power* of masked patches—the spectral quantity that clinical labels are defined by (Sec.). On the *same* criss-cross backbone, this one change removes *both* escapes at once—no codebook and no two-stage tokenizer (unlike LaBraM), and no reconstruction of task-irrelevant noise (unlike CBraMod/CSBrain)—yielding a single-stage objective aligned with what the tasks read out.

Latent-predictive SSL and collapse. Joint-embedding predictive architectures (I-JEPA (Assran et al. 2023)) predict latent representations of masked regions rather than pixels, and require an anti-collapse mechanism; LeJEPA (Balestriero and LeCun 2025) provides one (SIGReg) with provable isotropy. Our objective is not latent-predictive—its target is a fixed physical quantity (band power)—but it shares the collapse concern: we show that a *learnable* spectral target collapses to zero and that a *fixed* filterbank target removes the failure mode by construction. We retain SIGReg only as a light regularizer.

Data scale and benchmark leakage. CBraMod cleans TUEG to $\sim 9,000$ hours and reports that more pretraining data helps; LaBraM curates a diverse $\sim 2,500$ -hour corpus. We reach the opposite conclusion for spectral tasks: a ran-

dom $\sim 2,500$ -hour subset matches the full corpus at equal compute. Finally, TUEG-derived pretraining overlaps at the patient level with the in-corpus TUH downstream sets; we therefore evaluate only on benchmarks disjoint from TUEG for *every* method—sleep staging (ISRUC (Khalighi et al. 2016), HMC (Alvarez-Estevéz and Rijsman 2021)), depression (Mumtaz2016 (Mumtaz et al. 2018)) and mental arithmetic (Zyma et al. 2019)—using the subject-disjoint splits of CBraMod and CSBrain.

Method

Design principles. Three observations drive the design. (i) *The label lives in the spectrum*. Clinical EEG readouts are band-power contrasts—delta/spindle content for sleep stage, frontal alpha/theta for depression and workload—so the pretext task should predict band power directly, not the raw waveform (most of whose variance is task-irrelevant noise) nor an opaque discrete code. (ii) *Make the frequency decomposition explicit and learnable*. Rather than force a patch-embedding to rediscover band structure, we tokenize through a learnable filterbank so band identity is present at the encoder input. (iii) *Avoid a codebook*. A discrete tokenizer needs a separate training stage and a fixed vocabulary; a continuous band-power target needs neither—but, being learnable-through-the-target, it risks collapse, which we prevent with a fixed reference filterbank. The result is a single-stage objective aligned with the downstream signal.

Input. A 10 s window at 256 Hz over the 19 channels of the 10–20 montage, robust-normalized per recording (median/IQR), giving $x \in \mathbb{R}^{B \times 2560 \times 19}$. Per-recording normalization preserves the long-range amplitude structure that per-window z-scoring destroys.

Learnable filterbank tokenizer. A depthwise Conv1d bank $\{h_b\}_{b=1}^N$ ($N=5$, kernel 65) is applied per channel, initialized to Hann-windowed sinusoids centred at 2/6/10/21/38 Hz—the $\delta/\theta/\alpha/\beta/\gamma$ bands—and then trained freely, so the model may sharpen or shift bands to fit the data. Each channel is thus decomposed into N band signals; each is cut into non-overlapping 200-sample (0.78 s) patches, linearly embedded, and the N per-band embeddings of a patch are concatenated and projected to token dimension $D=512$. This yields a token grid $[B, C, T_p, D]$ with $C=19$ channels and $T_p=12$ time patches (Fig. 2).

Criss-cross encoder. The tokens are processed by a 12-layer transformer that, like CBraMod (Wang et al. 2025), alternates attention along the channel axis and the time axis within each layer (criss-cross), so spatial (montage) and temporal structure are modelled without the quadratic cost of full $C \times T_p$ attention. Learned positional embeddings are added on both axes.

Masking. We mask whole *time* patches (all bands of a masked patch), replacing their embeddings by a shared learnable mask token so the encoder never sees the hidden content. Masking whole patches rather than individual bands makes the pretext task predict a masked patch’s spectrum from its temporal context; an ablation (below) confirms time-patch masking outperforms band masking.

Spectral-prediction objective (primary). From the encoder output at each masked patch we predict the log band-power



Figure 1: **What should an EEG pretext task predict?** Prior foundation models predict a *proxy* of the signal: LaBraM predicts a discrete code from a separately trained vector-quantized tokenizer (a codebook plus two training stages); CBraMod and CSBrain reconstruct the *original raw waveform*, most of whose variance is task-irrelevant noise. We instead predict the *log band-power* of masked patches against a fixed filterbank target—the spectral quantity clinical labels are defined by. This single change removes *both* the codebook (unlike LaBraM) and the noise reconstruction (unlike CBraMod/CSBrain), in one single-stage objective.

of all N bands and supervise it against a target computed by a *fixed* (registered as a buffer, no gradient) reference filterbank h^{fix} applied to the raw signal:

$$\mathcal{L}_{\text{cf}} = \frac{1}{|\mathcal{M}|} \sum_{(c,t,b) \in \mathcal{M}} (\hat{p}_{c,t,b} - \log(1 + \bar{e}_{c,t,b}))^2, \quad \bar{e}_{c,t,b} = \frac{1}{P} \sum_{\tau} (h_b^{\text{fix}} * x)_{c,t,\tau}^2, \quad (1)$$

where $\mathcal{M}$ is the set of masked (channel, patch, band) triples and P the patch length. *Why a fixed target.* If the target filterbank were also learnable, the loss admits a trivial optimum: drive the target filters to zero so $\bar{e} \rightarrow 0$ and predict zero—we observe exactly this collapse ($\mathcal{L}_{\text{cf}} \rightarrow 4 \times 10^{-4}$ with degenerate, uninformative features). Freezing the target inverts the incentive: real band energy cannot be predicted from an uninformative encoding, so minimizing $\mathcal{L}_{\text{cf}}$ forces the encoder (and the learnable input filterbank) to stay informative. This is the crux of a codebook-free spectral objective and is verified in the ablations.

Regularization, and an inert reconstruction term. We add SIGReg (Balestriero and LeCun 2025), an isotropy penalty (weight $\lambda=0.05$) that discourages dimensional collapse, and use *no* latent JEPA term (an ablation found it inert). We also included a raw-signal reconstruction head (weight 0.5), but its loss is *identically zero* throughout pretraining: on the filterbank tokenizer the band tokens already determine the waveform, so reconstruction is trivially satisfied and contributes no gradient. The *effective* objective is therefore purely spectral,

$$\mathcal{L} = \mathcal{L}_{\text{cf}} + 0.05 \mathcal{L}_{\text{sig}}, \quad (2)$$

and Table 5 confirms the spectral term carries all of the performance.

Experiments

Pretraining. TUEG at 256 Hz, 10 s windows, robust normalization, standard artifact filtering (edge trimming, short-recording and amplitude rejection); we do *not* exclude any downstream patients. Our primary model uses a random $\sim 2,500$ -hour subset (2,100 patients drawn by a fixed seed); a matched-compute control uses the full $\sim 17,666$ hours (14,764 patients) with an identical gradient budget (same effective batch, $7\times$ fewer epochs).

Benchmarks (all leak-free). ISRUC, HMC: 5-class sleep staging (sequence-to-sequence). Mumtaz2016: depression (binary). Mental Arithmetic (eegmat): rest vs. arithmetic (binary). We use the CBraMod/CSBrain subject-disjoint splits.

Protocol and checkpoint policy. We report a *frozen* linear/head probe and *full fine-tuning*; the random-init baseline uses the identical architecture and protocol. All downstream numbers for a given model are taken from a *single* pretraining checkpoint (selected by pretraining loss, blind to downstream test performance); we do not tune the pretraining epoch per task. Because Mumtaz (11 test subjects) and Mental Arithmetic (4 test subjects) test sets are tiny, fine-tuning is high-variance on them; we therefore lead with the deterministic frozen probe and report fine-tuning as $\text{mean} \pm \text{std}$ over 5 seeds. Sleep sets are large; we report 3 reps.




Figure 2: **Codebook-free spectral-prediction pretraining.** A learnable filterbank tokenizer splits each patch into $\delta\theta\alpha\beta\gamma$ bands; time patches are masked out of the encoder input. From the encoder output the model regresses the log band-power of masked patches, supervised by a *fixed*, non-learnable filterbank target that cannot be collapsed; SIGReg adds an isotropy penalty. (A raw-signal reconstruction head we included is inert on this tokenizer—its loss stays zero—so the effective objective is purely spectral.)

Table 1: **ISRUC sleep staging (5-class).** Balanced accuracy, Cohen’s κ , weighted-F1 (mean $\pm$ std). Baselines as reported by Chen et al. (2025) on the identical subject-disjoint split; **bold** = best mean. Ours is pretrained on $\sim 2,500$ h vs. $\sim 9,000$ h (CBraMod/CSBrain).

Method	BA	κ	WF1
EEGNet	0.7154 $\pm$.0121	0.7040 $\pm$.0173	0.7513 $\pm$.0124
EEGConformer	0.7400 $\pm$.0133	0.7143 $\pm$.0162	0.7634 $\pm$.0151
SPaRCNet	0.7487 $\pm$.0075	0.7097 $\pm$.0132	0.7624 $\pm$.0092
ContraWR	0.7402 $\pm$.0126	0.7178 $\pm$.0156	0.7610 $\pm$.0137
CNN-Transformer	0.7363 $\pm$.0087	0.7129 $\pm$.0121	0.7719 $\pm$.0153
FFCL	0.7277 $\pm$.0182	0.7016 $\pm$.0292	0.7614 $\pm$.0197
ST-Transformer	0.7381 $\pm$.0205	0.7013 $\pm$.0352	0.7681 $\pm$.0175
BIOT	0.7527 $\pm$.0121	0.7291 $\pm$.0230	0.7790 $\pm$.0146
LaBraM-Base	0.7633 $\pm$.0102	0.7231 $\pm$.0182	0.7810 $\pm$.0133
CBraMod	0.7865 $\pm$.0110	0.7442 $\pm$.0152	0.8011 $\pm$.0099
CSBrain	0.7925 $\pm$.0030	0.7406 $\pm$.0102	0.7990 $\pm$.0091
Ours (FT)	0.7934 $\pm$.0051	0.7282 $\pm$.0010	0.7950 $\pm$.0012
Ours (frozen probe)	0.7522	0.6794	0.7503
Random init (frozen)	0.5657	0.4202	0.5235

Results

Draft: all numbers are the $\sim 2,500$ -hour model at its single converged checkpoint. Cells still marked ? are pending auxiliary runs (objective ablation, sleep frozen probe, random-init fine-tuning, full-corpus HMC control).

Sleep staging at one-seventh the data. Our $\sim 2,500$ -hour model reaches ISRUC balanced accuracy 0.7934 ± 0.0051 and HMC 0.7484 ± 0.0052 (Tables 1, 2), on par with CSBrain on ISRUC (0.7925) and *exceeding* it on HMC (0.7345; $+0.0139$, $\approx 2.7\sigma$), while CBraMod and CSBrain pretrain on $\sim 9,000$ hours. A full-data control on the $\sim 17,666$ -hour corpus—seven times more data at a matched gradient

Table 2: **HMC sleep staging (5-class).** Balanced accuracy, Cohen’s κ , weighted-F1 (mean $\pm$ std); baselines from Chen et al. (2025) on the identical split. **bold** = best mean.

Method	BA	κ	WF1
EEGNet	0.6534 $\pm$.0122	0.5886 $\pm$.0201	0.6536 $\pm$.0168
EEGConformer	0.7149 $\pm$.0086	0.6432 $\pm$.0055	0.7080 $\pm$.0039
SPaRCNet	0.4756 $\pm$.1109	0.3147 $\pm$.1315	0.4108 $\pm$.1310
ContraWR	0.4242 $\pm$.0541	0.2340 $\pm$.0554	0.2987 $\pm$.0288
CNN-Transformer	0.6573 $\pm$.0141	0.5961 $\pm$.0105	0.6896 $\pm$.0065
FFCL	0.4427 $\pm$.0702	0.2542 $\pm$.0654	0.2902 $\pm$.0485
ST-Transformer	0.2559 $\pm$.0141	0.0503 $\pm$.0183	0.1428 $\pm$.0122
BIOT	0.6862 $\pm$.0041	0.6295 $\pm$.0113	0.7091 $\pm$.0147
LaBraM-Base	0.7277 $\pm$.0101	0.6813 $\pm$.0053	0.7454 $\pm$.0027
CBraMod	0.7269 $\pm$.0041	0.6685 $\pm$.0104	0.7395 $\pm$.0089
CSBrain	0.7345 $\pm$.0047	0.6818 $\pm$.0046	0.7506 $\pm$.0042
Ours (FT)	0.7484 $\pm$.0052	0.6740 $\pm$.0160	0.7459 $\pm$.0131
Ours (frozen probe)	0.7099	0.6200	0.7030
Random init (frozen)	0.5003	0.3406	0.4528

budget—reaches ISRUC 0.7980, statistically indistinguishable from the $\sim 2,500$ -hour model (Table 4). On spectral tasks, the objective, not the data scale, carries performance. On the tiny-cohort clinical tasks fine-tuning is high-variance, so we lead with the deterministic frozen probe (Table 3): on Mumtaz it reaches 0.9173 BA / 0.9833 AUROC—*above* its own fine-tuned number (0.9129) and with zero seed variance—and on Mental Arithmetic 0.7083 BA / 0.7590 AUROC. The fine-tuned Mental Arithmetic mean (0.7625) even exceeds CSBrain’s 0.7558, but at ± 0.0727 we do not lean on it.

Checkpoint selection. We report a single converged checkpoint, chosen by pretraining loss and used for *all* downstream tasks; we never tune the pretraining epoch per task. It sits

Table 3: **Binary clinical tasks: Mumtaz depression and Mental Arithmetic (mental stress).** Balanced accuracy / AUC-PR / AUROC (mean $\pm$ std) on the Chen et al. (2025) subject-disjoint splits; baselines as reported there. **Bold** = best *robust* mean. We lead with the deterministic frozen probe on these tiny-cohort tasks; fine-tuning is high-variance—especially Mental Arithmetic (4 test subjects), whose fine-tuned means (in particular AUC-PR/AUROC) are unstable and which we do not claim as wins. [†]Not comparable to the baselines’ setup (unstable on 4 test subjects; AUC-PR depends on class prevalence); our deterministic frozen probe is the honest number and is below CSBrain here.

Method	Mumtaz (depression)			Mental Arithmetic		
	BA	AUC-PR	AUROC	BA	AUC-PR	AUROC
EEGNet	0.9232 $\pm$.0104	0.9626 $\pm$.0095	0.9639 $\pm$.0093	0.6770 $\pm$.0116	0.5763 $\pm$.0102	0.7321 $\pm$.0108
EEGConformer	0.9308 $\pm$.0117	0.9684 $\pm$.0105	0.9702 $\pm$.0101	0.6805 $\pm$.0123	0.5829 $\pm$.0134	0.7424 $\pm$.0128
SPaRCNet	0.9316 $\pm$.0095	0.9754 $\pm$.0065	0.9781 $\pm$.0063	0.6879 $\pm$.0107	0.5825 $\pm$.0193	0.7418 $\pm$.0132
ContraWR	0.9195 $\pm$.0115	0.9589 $\pm$.0102	0.9621 $\pm$.0092	0.6631 $\pm$.0097	0.5787 $\pm$.0164	0.7332 $\pm$.0082
CNN-Transformer	0.9305 $\pm$.0068	0.9757 $\pm$.0074	0.9742 $\pm$.0059	0.6779 $\pm$.0268	0.5777 $\pm$.0285	0.7258 $\pm$.0336
FFCL	0.9314 $\pm$.0083	0.9717 $\pm$.0021	0.9753 $\pm$.0033	0.6798 $\pm$.0142	0.5786 $\pm$.0266	0.7330 $\pm$.0198
ST-Transformer	0.9135 $\pm$.0103	0.9578 $\pm$.0086	0.9594 $\pm$.0059	0.6631 $\pm$.0173	0.5672 $\pm$.0259	0.7132 $\pm$.0174
BIOT	0.9358 $\pm$.0052	0.9736 $\pm$.0034	0.9758 $\pm$.0042	0.6875 $\pm$.0186	0.6004 $\pm$.0195	0.7536 $\pm$.0144
LaBraM-Base	0.9409 $\pm$.0079	0.9798 $\pm$.0093	0.9782 $\pm$.0057	0.6909 $\pm$.0125	0.5999 $\pm$.0155	0.7721 $\pm$.0093
CBraMod	0.9560 $\pm$.0056	0.9923 $\pm$.0032	0.9921 $\pm$.0025	0.7256 $\pm$.0132	0.6267 $\pm$.0099	0.7905 $\pm$.0073
CSBrain	0.9643 $\pm$.0155	0.9936 $\pm$.0034	0.9956 $\pm$.0020	0.7558 $\pm$.0106	0.6696 $\pm$.0221	0.8478 $\pm$.0297
Ours (frozen)	0.9173	0.9829	0.9833	0.7083	0.6239	0.7590
Ours (FT)	0.9129 $\pm$.0260	0.9817 $\pm$.0106	0.9809 $\pm$.0107	0.7625 $\pm$.0727	0.8045 $\pm$.0544 [†]	0.8995 $\pm$.0329 [†]

Table 4: **Data efficiency (ISRUC / HMC balanced accuracy).** Our $\sim$ 2,500-hour model vs. a matched-compute control on the full corpus and vs. prior methods with their pre-training budgets. Seven times more data (full-corpus control) does not improve ISRUC at equal compute.

Model	Pretrain data	ISRUC	HMC
LaBraM	$\sim$ 2,500 h	0.7633	0.7277
CBraMod	$\sim$ 9,000 h	0.7865	0.7269
CSBrain	$\sim$ 9,000 h	0.7925	0.7345
Ours (full ctrl)	$\sim$ 17,666 h	0.7980	?
Ours	$\sim$ 2,500 h	0.7934 $\pm$.0051	0.7484 $\pm$.0052

at roughly the matched gradient budget of the full-corpus control ($\sim 7\times$ more epochs on $\sim 7\times$ less data). Earlier checkpoints underfit; much later ones overfit the small subset: sleep accuracy still creeps up (ISRUC 0.7951, HMC 0.7522 at ep45) but the tiny-cohort clinical tasks degrade sharply (Mental Arithmetic falls from 0.7625 to 0.6174). Reporting every task at one converged checkpoint—rather than each at its own best epoch—is the honest choice, and coincides with the matched-compute comparison to the full-data control.

Frozen features already rival fine-tuning. A probe on our *frozen* encoder—no backbone update—reaches 0.7522 BA on ISRUC and 0.7099 on HMC (Tables 1, 2), matching fine-tuned BIOT (0.7527) and exceeding fine-tuned EEGNet (0.7154) and EEG Conformer (0.7400) on ISRUC, whereas the same architecture randomly initialized reaches only 0.5657/0.5003—a +0.19/+0.21 BA gap. The spectral

pretext task therefore produces a representation that is discriminative off-the-shelf, not only after fine-tuning.

Label efficiency. A probe on our frozen features reaches high accuracy with far fewer downstream labels than a random encoder (Fig. 3). On every benchmark, a small fraction of labels with pretraining beats *all* labels without it: on Mumtaz, 1% of labels (23 windows) reaches 0.8990 BA, above a random encoder using 100% (0.8200); on ISRUC, 10% of sequences reaches 0.6520, above random at 100% (0.5540); similarly for HMC (0.5560 at 10% vs. 0.4920 at 100%) and Mental Arithmetic (0.5790 at 1% vs. 0.5160 at 100%). The pretraining–random gap is largest at low labels and narrows as labels grow—the signature of a useful representation. On the sequence tasks the gap emerges at 10%: below that, too few sequences exist to fit the temporal head at all (both near chance), so pretraining lowers the label threshold at which the task becomes learnable by an order of magnitude.

Ablations

Spectral prediction is the load-bearing objective. Table 5 isolates the two loss terms at identical data and compute. Removing the reconstruction term (CF-only) barely changes anything—ISRUC 0.7864, HMC 0.7486, within noise of the full config (0.7934/0.7484)—as expected, since that term’s loss is identically zero and it contributes no gradient (Sec.). Removing the spectral term instead (MAE-only) leaves no informative objective and is catastrophic: ISRUC and HMC collapse to 0.2124 and 0.2737, the 5-class chance level ($\kappa \approx 0.05$), *below* even a randomly initialized encoder (ISRUC 0.738 fine-tuned). On our filterbank backbone the band tokens already determine the waveform, so raw-signal recon-




Figure 3: **Label efficiency** of the $\sim 2,500$ -hour model. Balanced accuracy of a probe trained on a fraction of downstream labels, pretrained (ours, solid) vs. random init (dashed); dotted line is random at 100% labels, shaded band is the pretrained–random gap. A small fraction of labels with pretraining exceeds all labels without it on every task; on the sequence tasks (ISRUC, HMC) the advantage appears once $\geq 10\%$ of sequences suffice to fit the head.

Table 5: **Objective ablation** at $\sim 2,500$ h, matched compute (same subset, same steps; only the loss changes). Isolates spectral band-power prediction (CF) vs. raw-signal reconstruction (MAE). Reported on the low-variance sleep benchmarks (large test sets, std ~ 0.005), where a small delta is resolvable; the tiny-cohort clinical tasks are too high-variance to separate the conditions.

Objective	ISRUC BA	HMC BA
MAE only (reconstruction)	$0.2124 \pm .0015$	$0.2737 \pm .0247$
CF only (spectral prediction)	$0.7864 \pm .0023$	$0.7486 \pm .0017$
CF + MAE (ours)	$0.7934 \pm .0051$	$0.7484 \pm .0052$

struction is trivially satisfied and teaches nothing; only the lossy, structure-selecting band-power target compels the encoder to learn. Band-power prediction is thus both necessary and sufficient—the sole driver of the representation.

Fixed vs. learnable target. A learnable-target filterbank collapses ($\mathcal{L}_{cf} \rightarrow 0$, degenerate features); the fixed target is necessary and sufficient to prevent this.

Limitations

Our model is designed for the spectral, quasi-stationary structure of clinical and passive EEG. It is *not* suited to motor imagery, whose discriminative signal is a spatially localized event-related (de)synchronization rather than a stationary band-power pattern; we exclude motor-imagery benchmarks by design rather than reporting weak numbers. Clinical

test cohorts are small (Mumtaz 11, Mental Arithmetic 4 test subjects), so fine-tuning estimates are high-variance and we lead with the deterministic frozen probe. We also exclude TUH-derived abnormality detection (TUAB) from the main comparison: it overlaps at the patient level with TUEG pre-training for every method, and is off-topic for our spectral scope.

Conclusion

Predicting band power—directly, against a fixed filterbank target, with no codebook and no two-stage tokenizer—is a simple objective that is well matched to the spectral structure of clinical EEG. It reaches state-of-the-art sleep staging with a fraction of the pretraining data used by prior foundation models, is label-efficient downstream, and makes the surprising point that for spectral EEG tasks the pretext *objective*, not the data scale, is what matters. We argue this reframes what an EEG foundation model should predict.

References

- Alvarez-Estevez, D.; and Rijsman, R. M. 2021. Inter-Database Validation of a Deep Learning Approach for Automatic Sleep Scoring. *PLOS ONE*, 16(8): e0256111.
- Assran, M.; Duval, Q.; Misra, I.; Bojanowski, P.; Vincent, P.; Rabbat, M.; LeCun, Y.; and Ballas, N. 2023. Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.

- Balestriero, R.; and LeCun, Y. 2025. LeJEPa: Provable and Scalable Self-Supervised Learning Without the Heuristics. *arXiv preprint arXiv:2511.08544*.
- Bao, H.; Dong, L.; Piao, S.; and Wei, F. 2022. BEiT: BERT Pre-Training of Image Transformers. In *International Conference on Learning Representations (ICLR)*.
- Chen, Y.; et al. 2025. CSBrain: A Cross-Scale Spatiotemporal Brain Foundation Model for EEG Decoding. *arXiv preprint arXiv:2506.00000*. Author list to be completed from the published version.
- He, K.; Chen, X.; Xie, S.; Li, Y.; Dollár, P.; and Girshick, R. 2022. Masked Autoencoders Are Scalable Vision Learners. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Jiang, W.; Zhao, L.; and Lu, B.-L. 2024. Large Brain Model for Learning Generic Representations with Tremendous EEG Data in BCI. In *International Conference on Learning Representations (ICLR)*.
- Jing, J.; et al. 2023. Development of Expert-Level Classification of Seizures and Rhythmic and Periodic Patterns During EEG Interpretation. *Neurology*, 100(17): e1750–e1762.
- Khalighi, S.; Sousa, T.; Santos, J. M.; and Nunes, U. 2016. ISRUC-Sleep: A Comprehensive Public Dataset for Sleep Researchers. *Computer Methods and Programs in Biomedicine*, 124: 180–192.
- Lawhern, V. J.; Solon, A. J.; Waytowich, N. R.; Gordon, S. M.; Hung, C. P.; and Lance, B. J. 2018. EEGNet: A Compact Convolutional Neural Network for EEG-based Brain–Computer Interfaces. *Journal of Neural Engineering*, 15(5): 056013.
- Li, H.; et al. 2022. A Feature Fusion and Contrastive Learning Framework for EEG Decoding. *arXiv preprint arXiv:2209.00000*. Bibliographic details to be verified.
- Mumtaz, W.; Ali, S. S. A.; Yasin, M. A. M.; and Malik, A. S. 2018. A Machine Learning Framework Involving EEG-Based Functional Connectivity to Diagnose Major Depressive Disorder (MDD). *Medical & Biological Engineering & Computing*, 56: 233–246.
- Obeid, I.; and Picone, J. 2016. The Temple University Hospital EEG Data Corpus. *Frontiers in Neuroscience*, 10: 196.
- Peh, W. Y.; Yao, Y.; and Dauwels, J. 2022. Transformer Convolutional Neural Networks for Automated Artifact Detection in Scalp EEG. In *44th Annual International Conference of the IEEE Engineering in Medicine & Biology Society (EMBC)*.
- Song, Y.; Jia, X.; Yang, L.; and Xie, L. 2021. Transformer-Based Spatial-Temporal Feature Learning for EEG Decoding. *arXiv preprint arXiv:2106.11170*.
- Song, Y.; Zheng, Q.; Liu, B.; and Gao, X. 2023. EEG Conformer: Convolutional Transformer for EEG Decoding and Visualization. *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, 31: 710–719.
- van den Oord, A.; Vinyals, O.; and Kavukcuoglu, K. 2017. Neural Discrete Representation Learning. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Wang, J.; Zhao, S.; Luo, Z.; Zhou, Y.; Jiang, H.; Li, S.; Li, T.; and Pan, G. 2025. CBraMod: A Criss-Cross Brain Foundation Model for EEG Decoding. In *International Conference on Learning Representations (ICLR)*.
- Yang, C.; Westover, M. B.; and Sun, J. 2023. BIOT: Biosignal Transformer for Cross-Data Learning in the Wild. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Yang, C.; Xiao, C.; Westover, M. B.; and Sun, J. 2023. Self-Supervised EEG Representation Learning for Automatic Sleep Staging. *JMIR AI*, 2: e46769.
- Zyma, I.; Tukaev, S.; Seleznev, I.; Kiyono, K.; Popov, A.; Chernykh, M.; and Shpenkov, O. 2019. Electroencephalograms during Mental Arithmetic Task Performance. *Data*, 4(1): 14.